

# **Computer Vision**

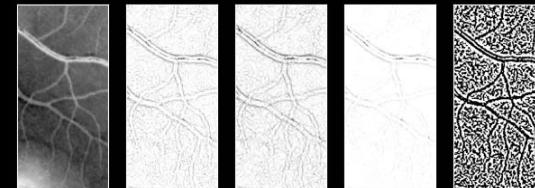
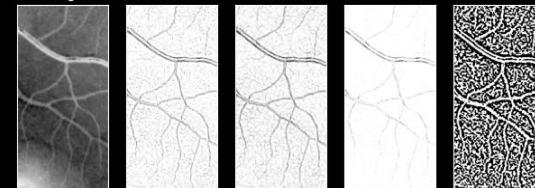
## *Hackathon Crash Course*

Sam Sweere  
Werner van der Veen

*Purpose: gain high-level understanding from digital images or videos*

Approaches:

- Classical
  - Fast & lightweight, but limited
  - Deterministic
  - Easily parametrizable
  - Suitable for edge computing, baselines
  - Examples: edge detection, denoising, inpainting, thresholding → *features*
- Deep learning
  - Heavy but powerful
  - Compute/data hungry
  - Can also be interpretable
  - Examples: classification, detection, segmentation, generation → *meaning*

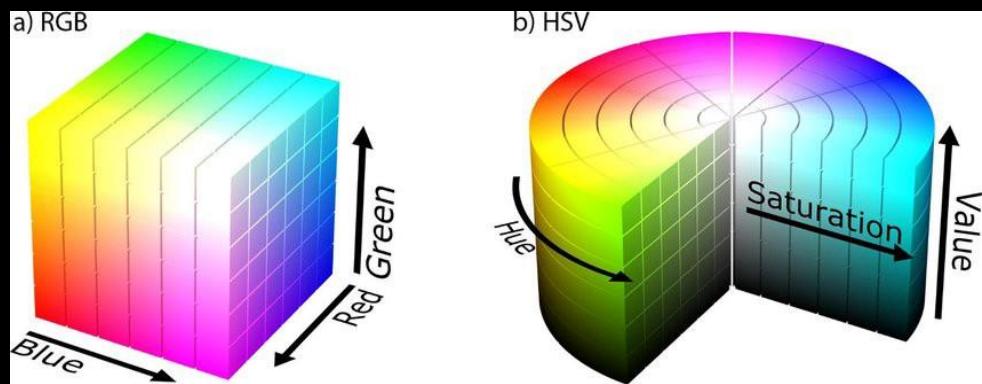


## *Classical Computer Vision methods*

— Before we start: Handling image data —

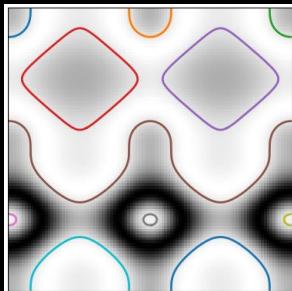
- Images are rasters of pixel values ( $W \times H \times C$ )
  - np.ndarray
  - (R, G, B) vs. (H, S, V)
  - opencv, PIL, imageio
  - Often np.bool, np.uint8, np.uint16 or np.float16
  - Easily displayed using plt.imshow()
  - scikit-image

```
import imageio  
  
im = imageio.imread("image.png")  
  
print(im.shape)  
>>> (100, 200, 3)
```



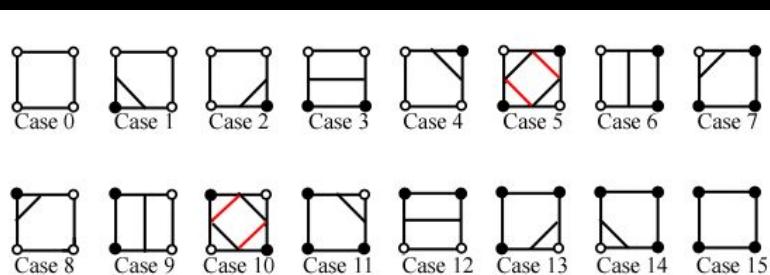
## *Classical Computer Vision methods*

— Edges: finding iso-valued contours —



### **Marching Squares:**

```
skimage.measure.find_contours()
```



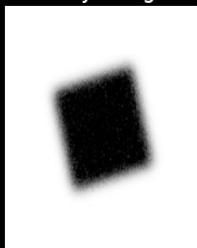
### Algorithm:

In every grid of 2x2 pixels, draw an isoline according to whether its 4 pixels are lower or higher than the isovalue.

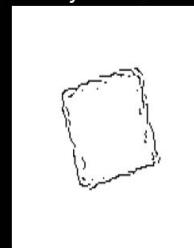
## *Classical Computer Vision methods*

### — Edges: edge detection —

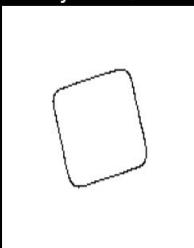
noisy image



Canny filter,  $\sigma = 1$



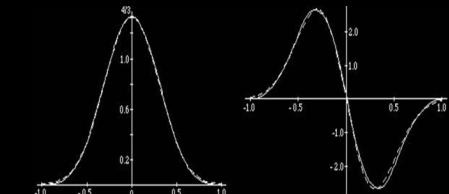
Canny filter,  $\sigma = 3$



Canny filter

Multi-stage:

1. derivative of Gaussian (parameter  $\sigma$  is width)
2. edges thinned to 1-pixel by keeping only gradient maxima
3. remove weak, unconnected edges



Sobel filter

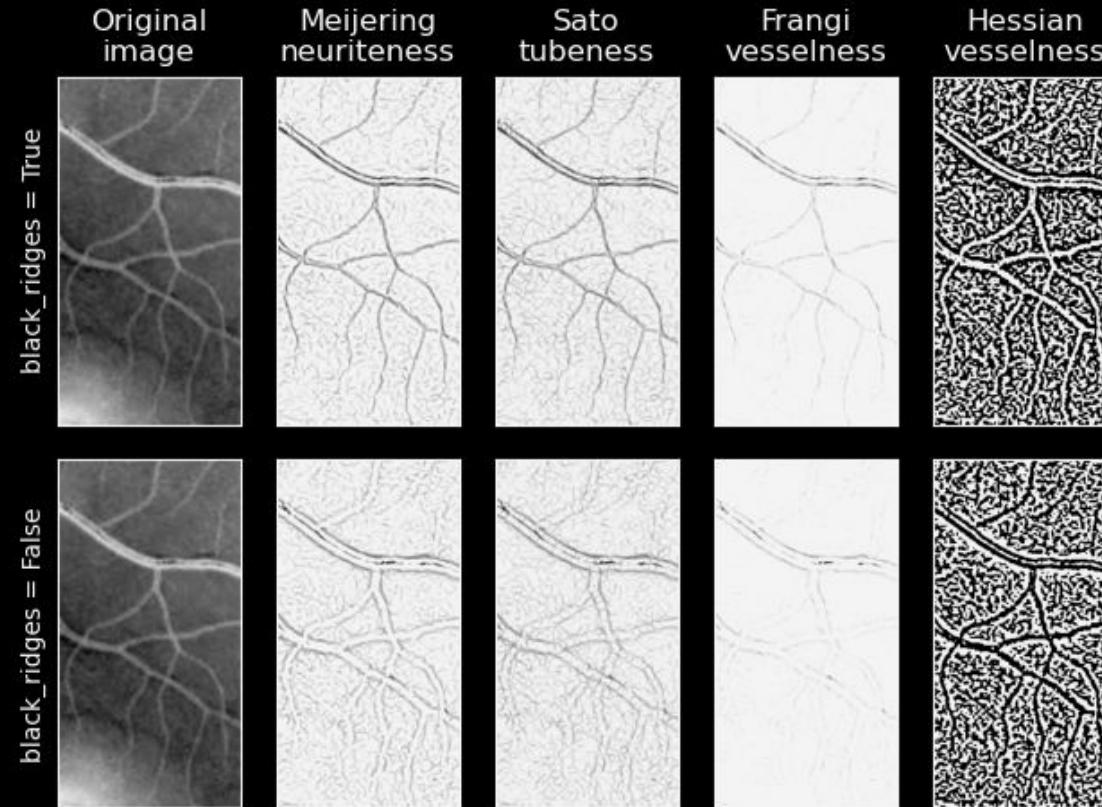


$$\mathbf{G}_x = \begin{bmatrix} -1 & 0 & +1 \\ -2 & 0 & +2 \\ -1 & 0 & +1 \end{bmatrix} * \mathbf{A} \quad \text{and} \quad \mathbf{G}_y = \begin{bmatrix} +1 & +2 & +1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{bmatrix} * \mathbf{A}$$



$$\mathbf{G} = \sqrt{\mathbf{G}_x^2 + \mathbf{G}_y^2} \quad \Theta = \text{atan2}(\mathbf{G}_y, \mathbf{G}_x)$$

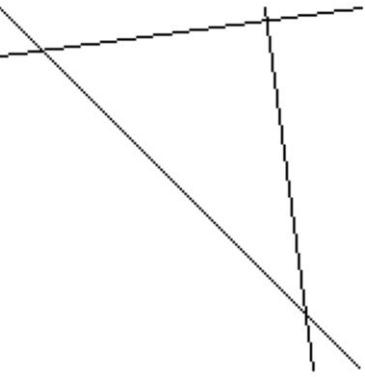
*Classical Computer Vision methods*  
— Edges: ridge filters —



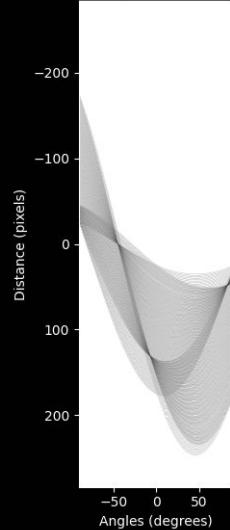
## *Classical Computer Vision methods*

### — Straight line detection: Hough transform —

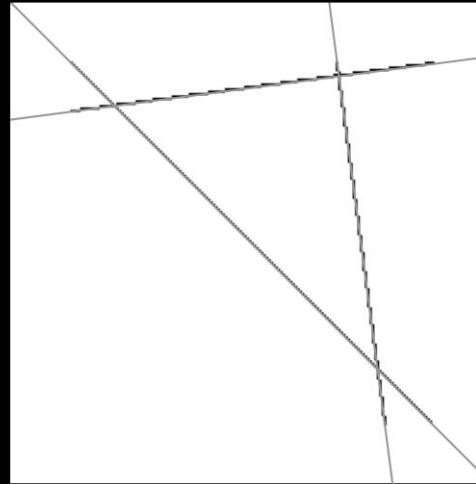
Input image



Hough transform



Detected lines



Algorithm:

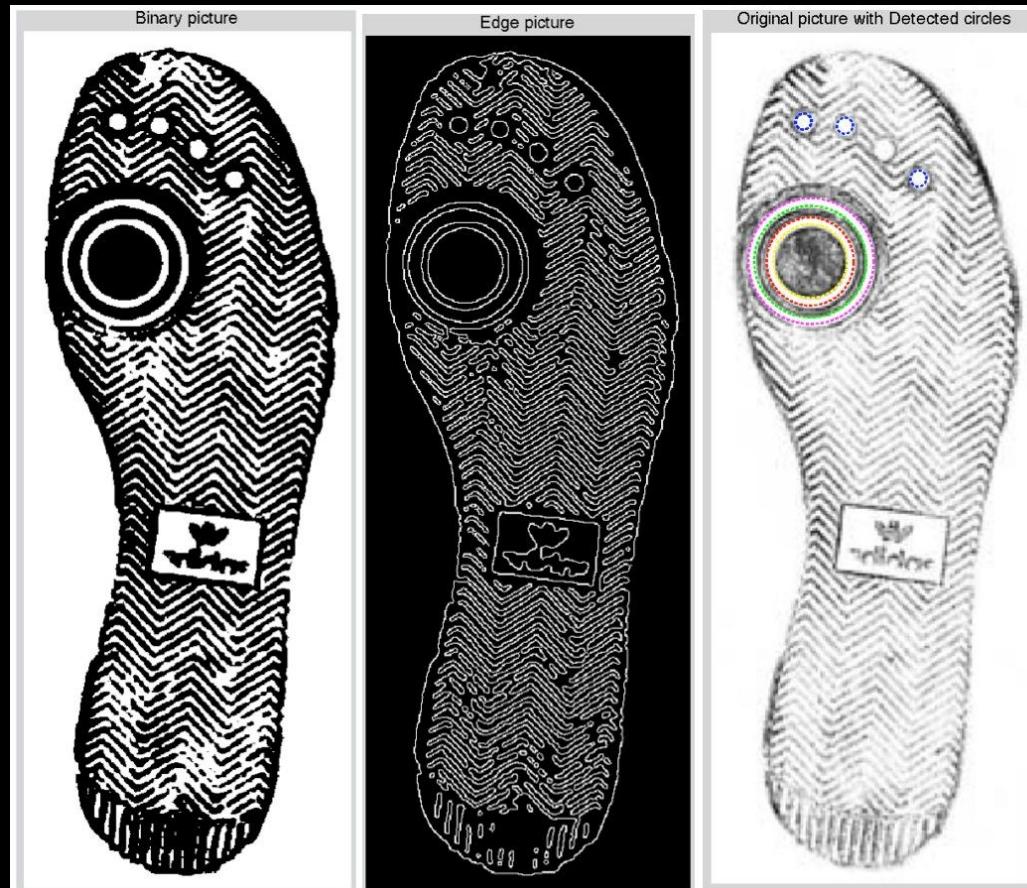
1. For each point, draw lines with angle  $\theta$  through it and find length  $r$  of its vector orthogonal to the origin.
2. Draw all  $(\theta, r)$  pairs of all points in heatmap
3. Local heatmap maxima indicate lines in image

## *Classical Computer Vision methods*

### — Circle detection: Hough transform —

Algorithm:

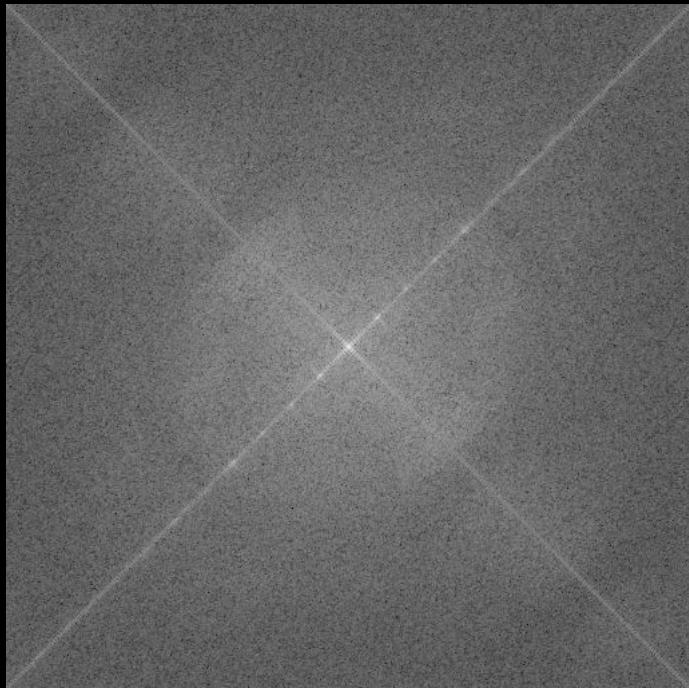
1. For each pixel  $(x, y)$  and every radius  $r \in R$ , count the number of edge pixels on that circle.
2. Draw all  $(x, y, r)$  pairs of all points in a 3D heatmap.
3. Local heatmap maxima indicate circles in the original image



## *Classical Computer Vision methods*

### — Corner detection: Fourier transform —

- Decompose image signal into spectral components
- Analysis of
  - speed of change (distance from center)
  - direction (angle from center)
- Can be combined with e.g. edge detection



## *Classical Computer Vision methods* — Corner detection: Moravec transform —

Algorithm:

1. For all local patches, the center is a corner if the patch is not similar to *any* nearby patches.
2. Corner strength is the dissimilarity.

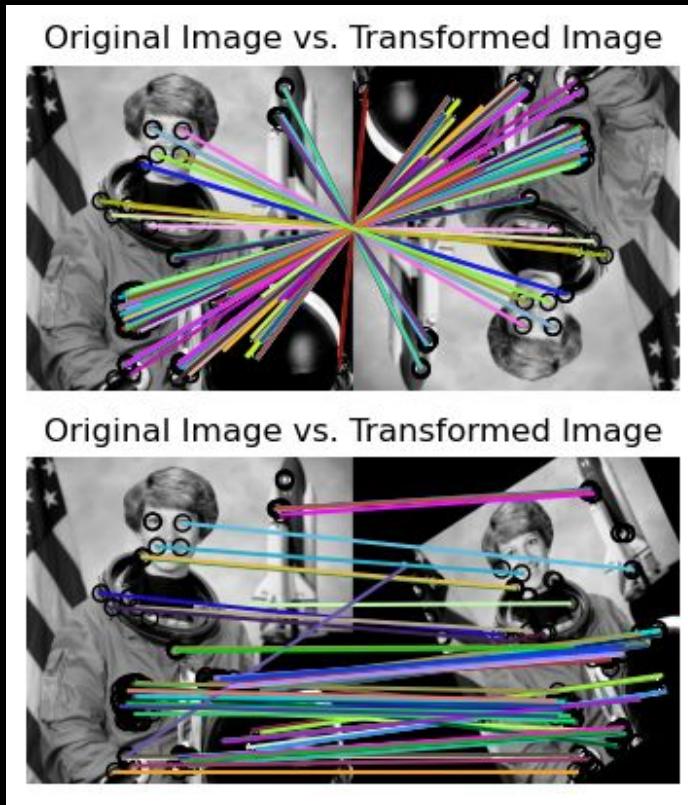


## *Classical Computer Vision methods*

### — Matching: ORB —

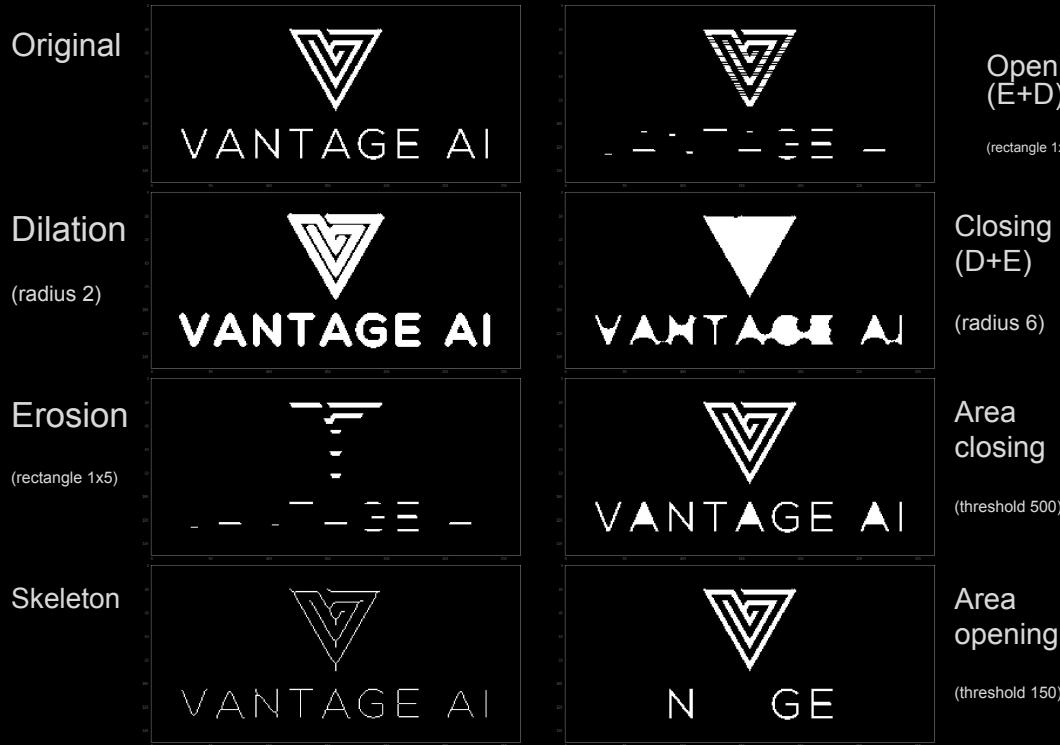
#### Oriented FAST and rotated BRIEF (ORB)

- Comparatively rotation & scale invariant
- FAST: efficient corner detection method
- BRIEF: “fingerprints” keypoints for matching
- Fast enough for real-time applications



## *Classical Computer Vision methods*

### — Morphological processing —



Opening  
(E+D)

(rectangle 1x4)

Closing  
(D+E)

(radius 6)

Area  
closing

(threshold 500)

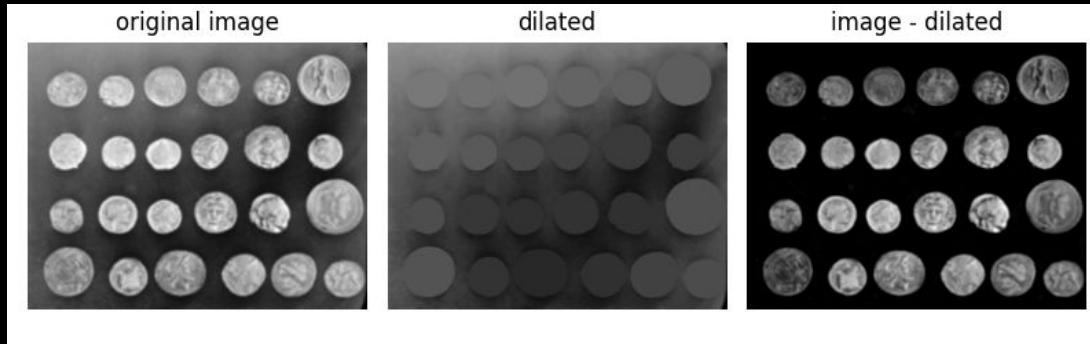
Area  
opening

(threshold 150)

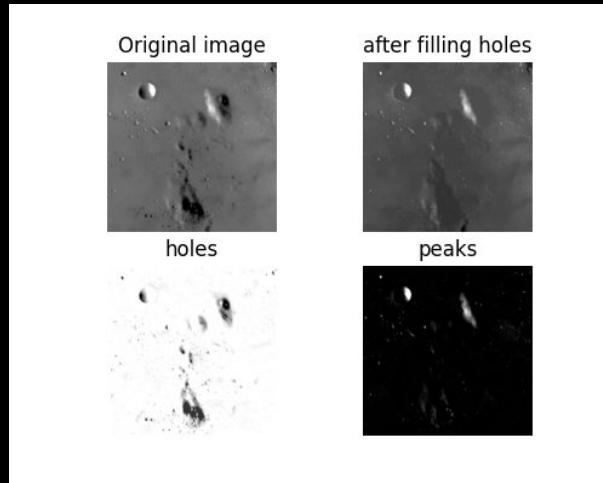
## *Classical Computer Vision methods*

### *— Stacking morphological processing: simple and effective —*

Reconstruction  
by dilation



Reconstruction by erosion  
and dilation



Method: use “seed” to specify  
values that spread, and “mask”  
to clip the values.

## *Classical Computer Vision methods* — Inpainting: biharmonic equation —

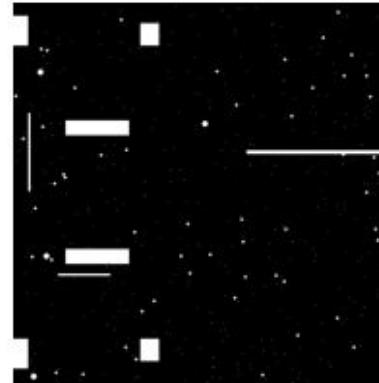
Biharmonic equation:

- Uses 4th-order partial differential equation to interpolate between surrounding patterns

Original image



Mask



Defected image



Inpainted image



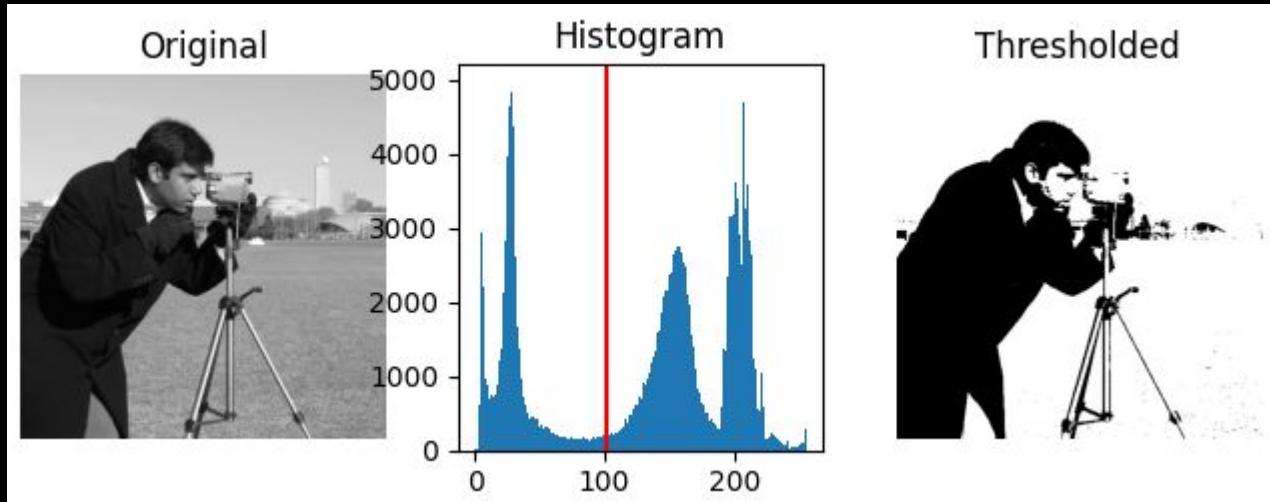
# *Classical Computer Vision methods*

## — Thresholding —

<p>Original</p> <p><b>Region-based segmentation</b></p> <p>Let us first determine markers of the coins and the background. These markers are pixels that we can label unambiguously as either object or background. Here, the markers are found at the two extreme parts of the histogram of grey values:</p> <pre>&gt;&gt;&gt; markers = np.zeros_like(coins)</pre> <p>Li</p> <p><b>Region-based segmentation</b></p> <p>Let us first determine markers of the coins and the background. These markers are pixels that we can label unambiguously as either object or background. Here, the markers are found at the two extreme parts of the histogram of grey values:</p> <pre>&gt;&gt;&gt; markers = np.zeros_like(coins)</pre> <p>Minimum</p> <p><b>Region-based segmentation</b></p> <p>Let us first determine markers of the coins and the background. These markers are pixels that we can label unambiguously as either object or background. Here, the markers are found at the two extreme parts of the histogram of grey values:</p> <pre>&gt;&gt;&gt; markers = np.zeros_like(coins)</pre> <p>Triangle</p> <p><b>Region-based segmentation</b></p> <p>Let us first determine markers of the coins and the background. These markers are pixels that we can label unambiguously as either object or background. Here, the markers are found at the two extreme parts of the histogram of grey values:</p> <pre>&gt;&gt;&gt; markers = np.zeros_like(coins)</pre>	<p>Isodata</p> <p><b>Region-based segmentation</b></p> <p>Let us first determine markers of the coins and the background. These markers are pixels that we can label unambiguously as either object or background. Here, the markers are found at the two extreme parts of the histogram of grey values:</p> <pre>&gt;&gt;&gt; markers = np.zeros_like(coins)</pre> <p>Mean</p> <p><b>Region-based segmentation</b></p> <p>Let us first determine markers of the coins and the background. These markers are pixels that we can label unambiguously as either object or background. Here, the markers are found at the two extreme parts of the histogram of grey values:</p> <pre>&gt;&gt;&gt; markers = np.zeros_like(coins)</pre> <p>Otsu</p> <p><b>Region-based segmentation</b></p> <p>Let us first determine markers of the coins and the background. These markers are pixels that we can label unambiguously as either object or background. Here, the markers are found at the two extreme parts of the histogram of grey values:</p> <pre>&gt;&gt;&gt; markers = np.zeros_like(coins)</pre> <p>Yen</p> <p><b>Region-based segmentation</b></p> <p>Let us first determine markers of the coins and the background. These markers are pixels that we can label unambiguously as either object or background. Here, the markers are found at the two extreme parts of the histogram of grey values:</p> <pre>&gt;&gt;&gt; markers = np.zeros_like(coins)</pre>
--	---

## *Classical Computer Vision methods*

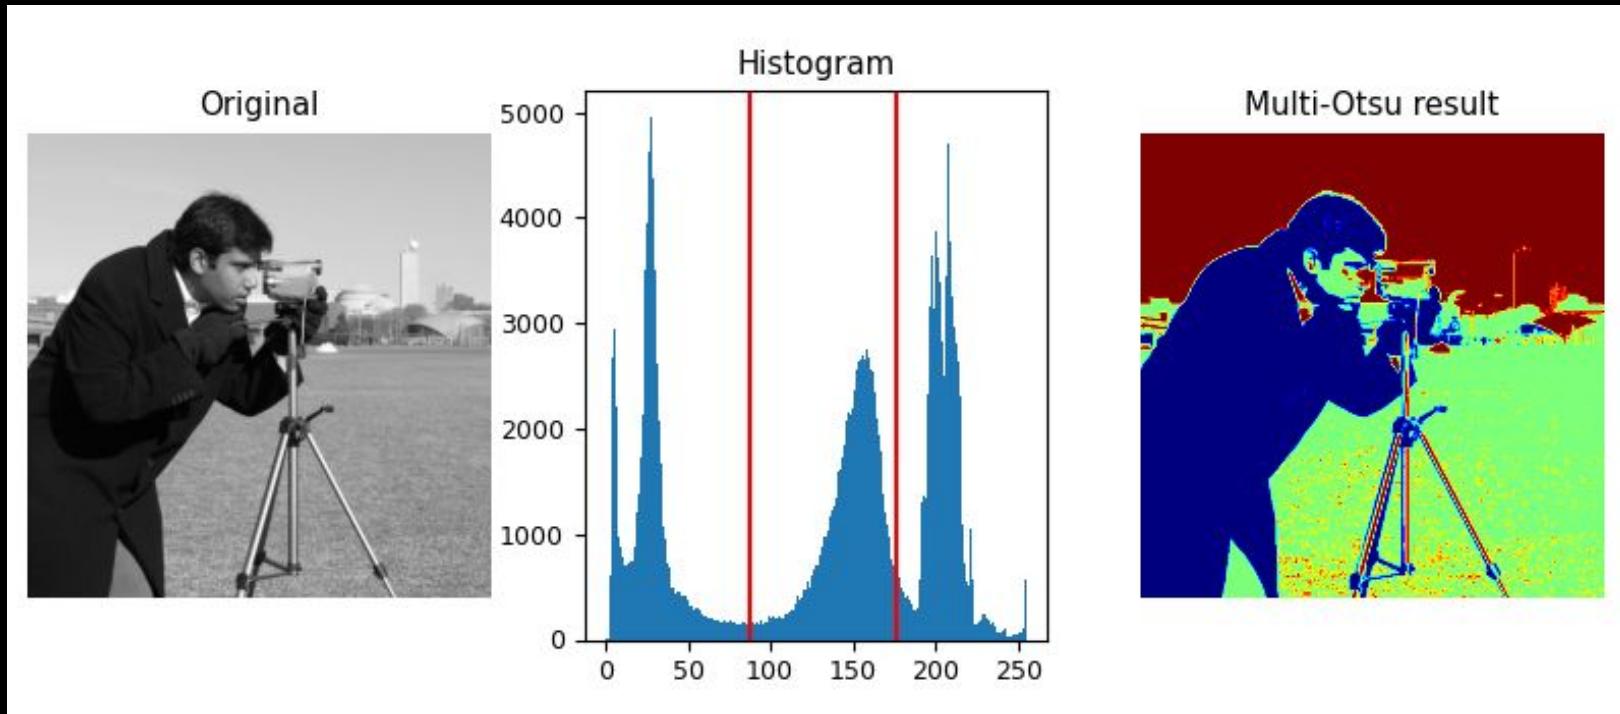
### *— Thresholding: Otsu—*



“Optimal” threshold:

- Maximum variance between two classes
- (and therefore minimal intra-class variance)

*Classical Computer Vision methods*  
— Thresholding: multi-Otsu —



## *Classical Computer Vision methods*

### — Thresholding: local Otsu —

#### Global thresholding

##### Region-based segmentation

Let us first determine markers of the coins and the background. These markers are pixels that we can label unambiguously as either object or background. Here, the markers are found at the two extreme parts of the histogram of grey values:

```
>>> markers = np.zeros_like(coins)
```

- Much slower, but more accurate
- Local neighborhood

#### Local thresholding

##### Region-based segmentation

Let us first determine markers of the coins and the background. These markers are pixels that we can label unambiguously as either object or background. Here, the markers are found at the two extreme parts of the histogram of grey values:

```
>>> markers = np.zeros_like(coins)
```

#### Original

##### Region-based segmentation

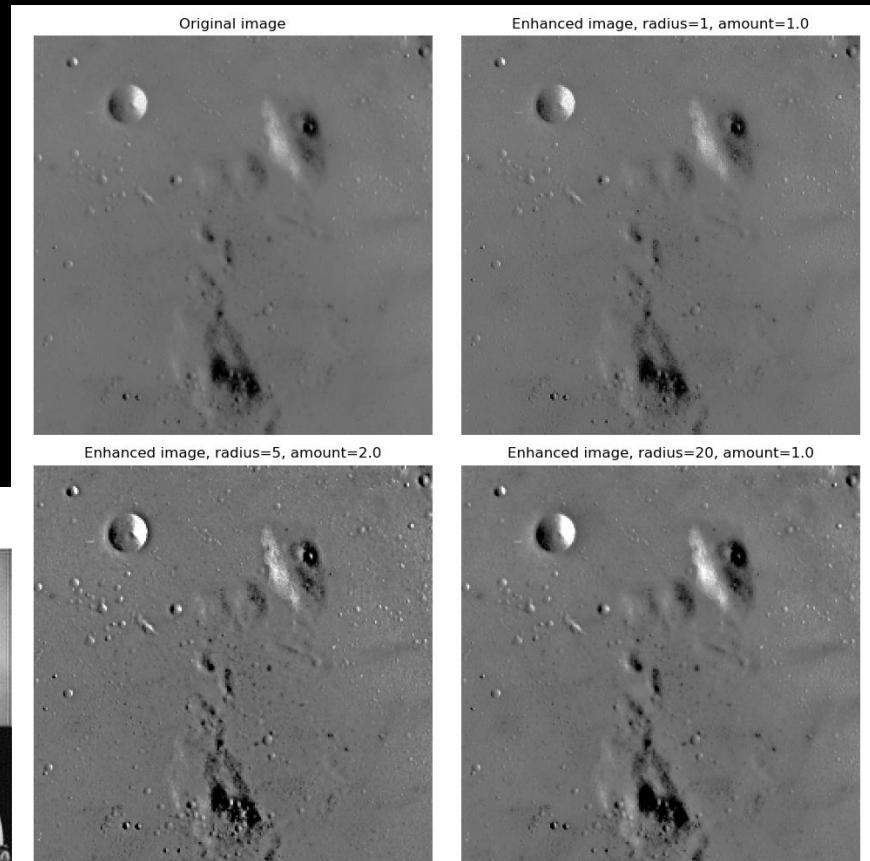
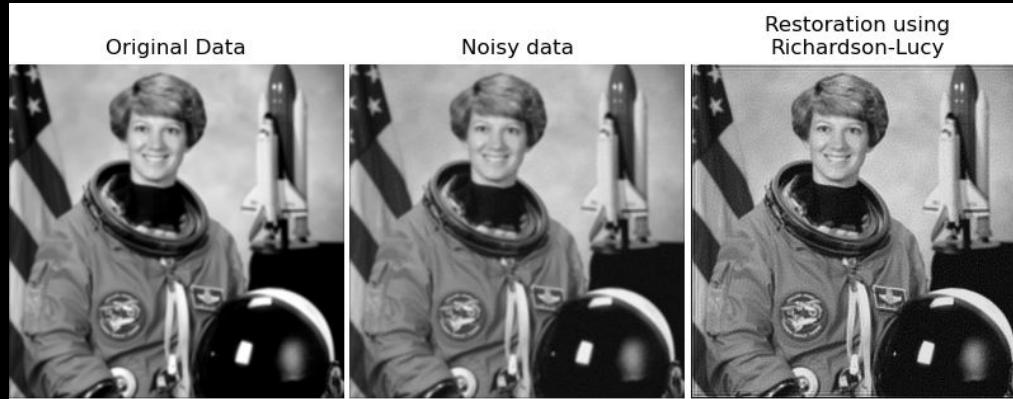
Let us first determine markers of the coins and the background. These markers are pixels that we can label unambiguously as either object or background. Here, the markers are found at the two extreme parts of the histogram of grey values:

```
>>> markers = np.zeros_like(coins)
```

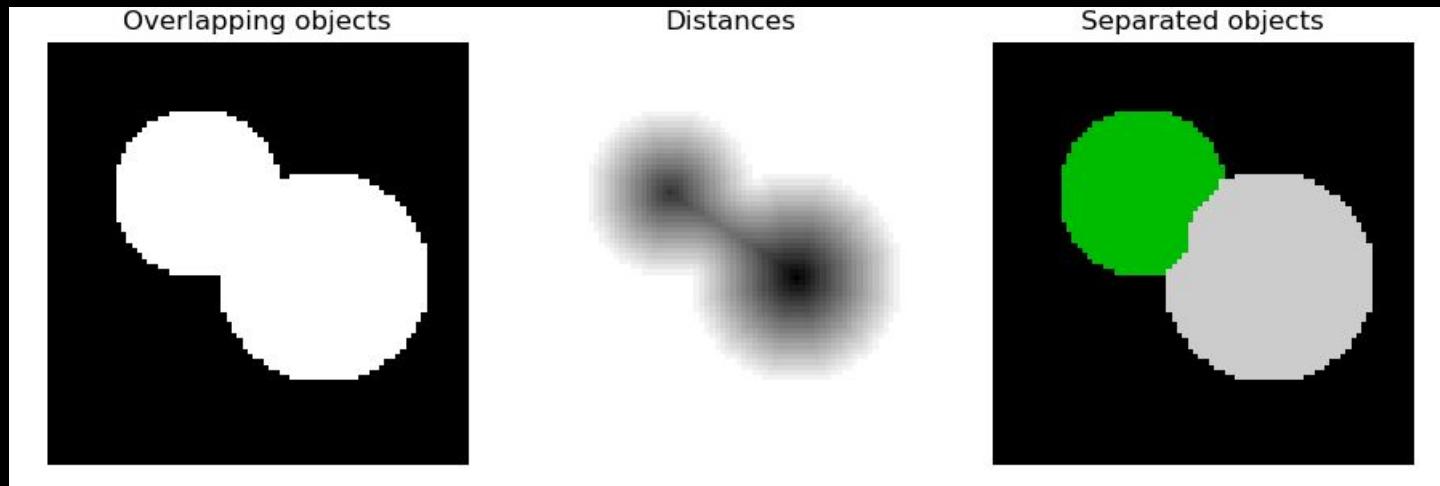
## *Classical Computer Vision methods*

### — Sharpening —

- Difference between original and blurred, scaled back, and added to original



*Classical Computer Vision methods*  
— Segmentation: Watershed algorithm —



1. Calculate distance from background
2. Flood local maxima in distance map
3. Edges between basins are object boundaries

## *Deep Learning Computer Vision methods*

- Classification
- Object detection
- Segmentation
- Generation

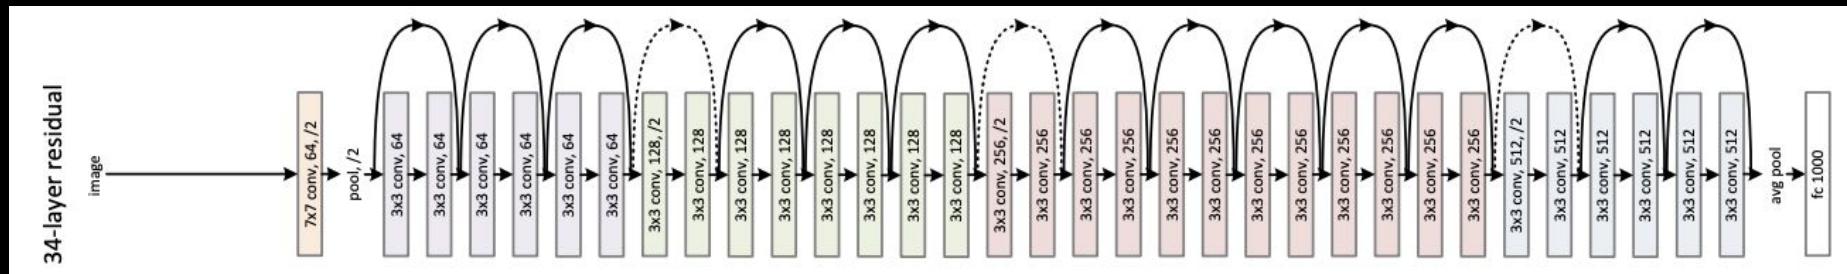
# *Deep Learning Computer Vision methods*

## — Classification —

Pretty much always CNNs

Tried and true: ResNet pretrained on ImageNet

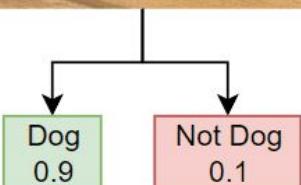
- Transfer learning (unfreeze only last  $N$  layers)
- Rescale input to  $[0, 1]$
- Data augmentation & feature engineering
- Saliency maps
  - (Gradient weighted) class activation mapping



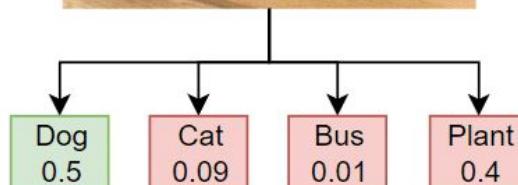
# *Deep Learning Computer Vision methods*

## — Classification —

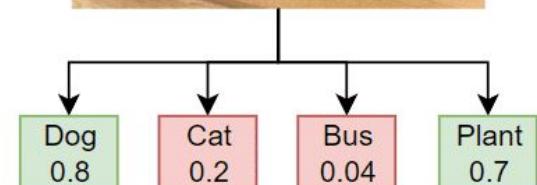
**Binary Classification**



**Multiclass Classification**



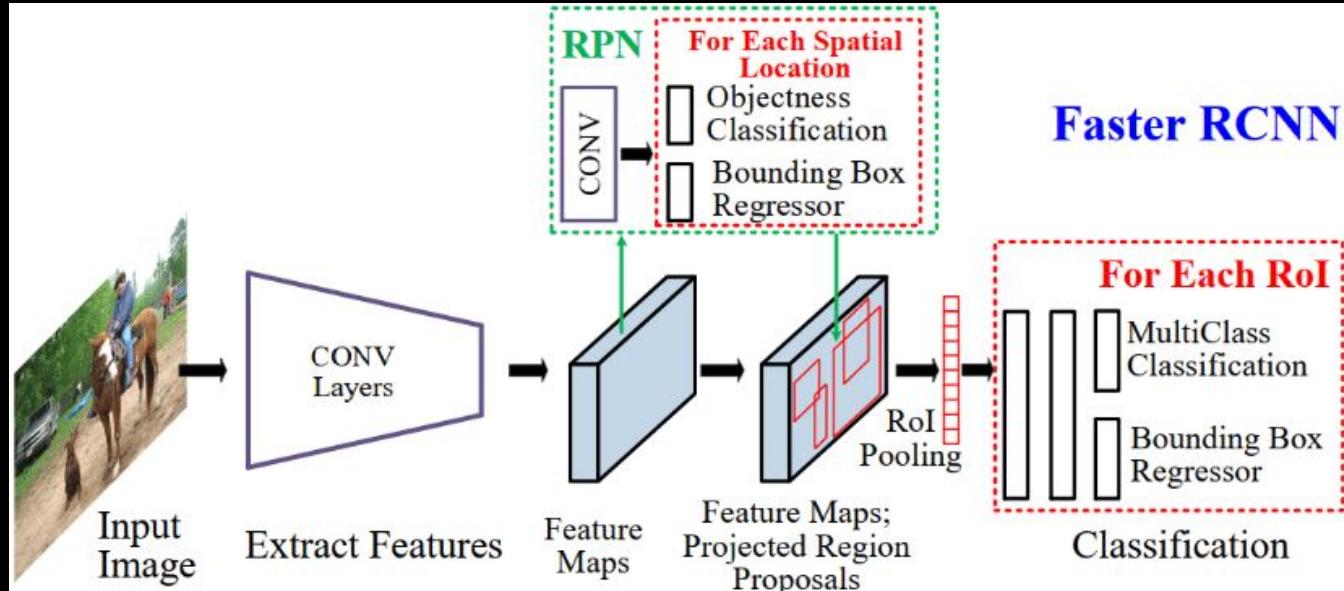
**Multilabel Classification**



## *Deep Learning Computer Vision methods*

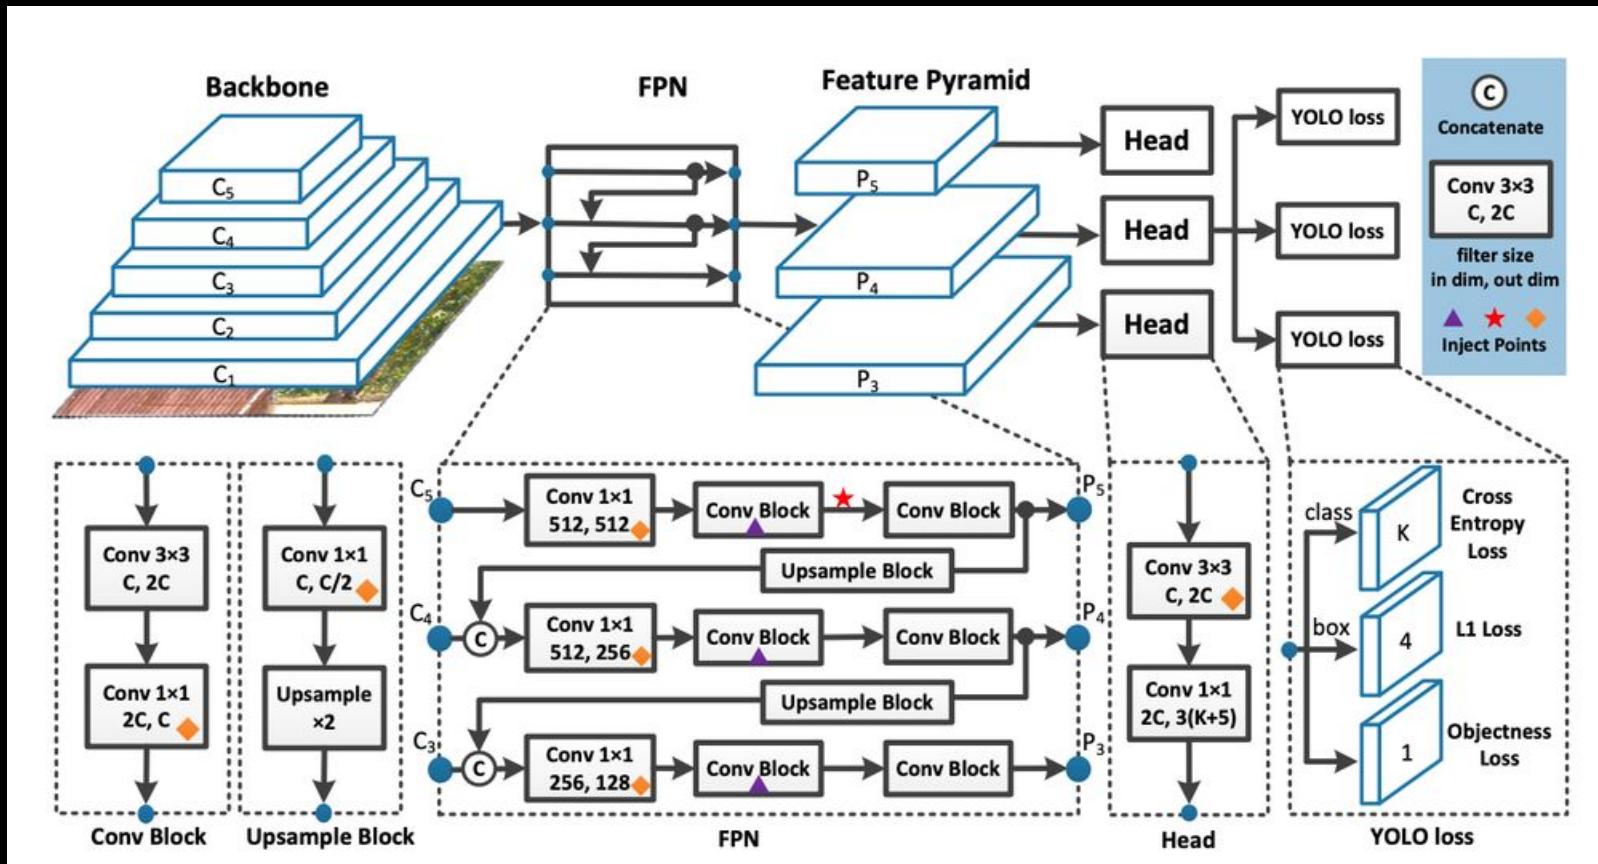
— Object detection: two-shot (Faster RCNN) —

1. Feature extraction
2. Anchor box proposals
3. Prediction heads
4. Regression heads

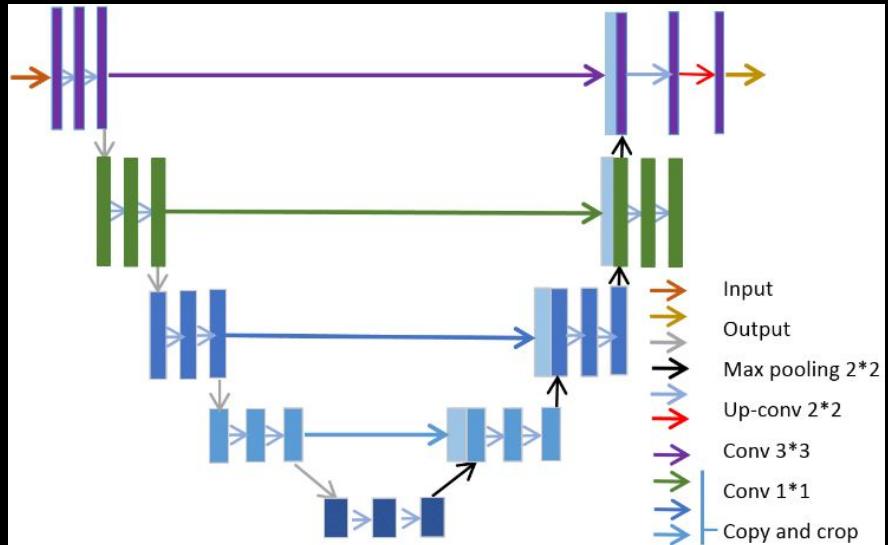


# Deep Learning Computer Vision methods

## — Object detection: one-shot (YOLOv7) —



## *Deep Learning Computer Vision methods* — Segmentation: U-Net (2014) —



Semantic segmentation:

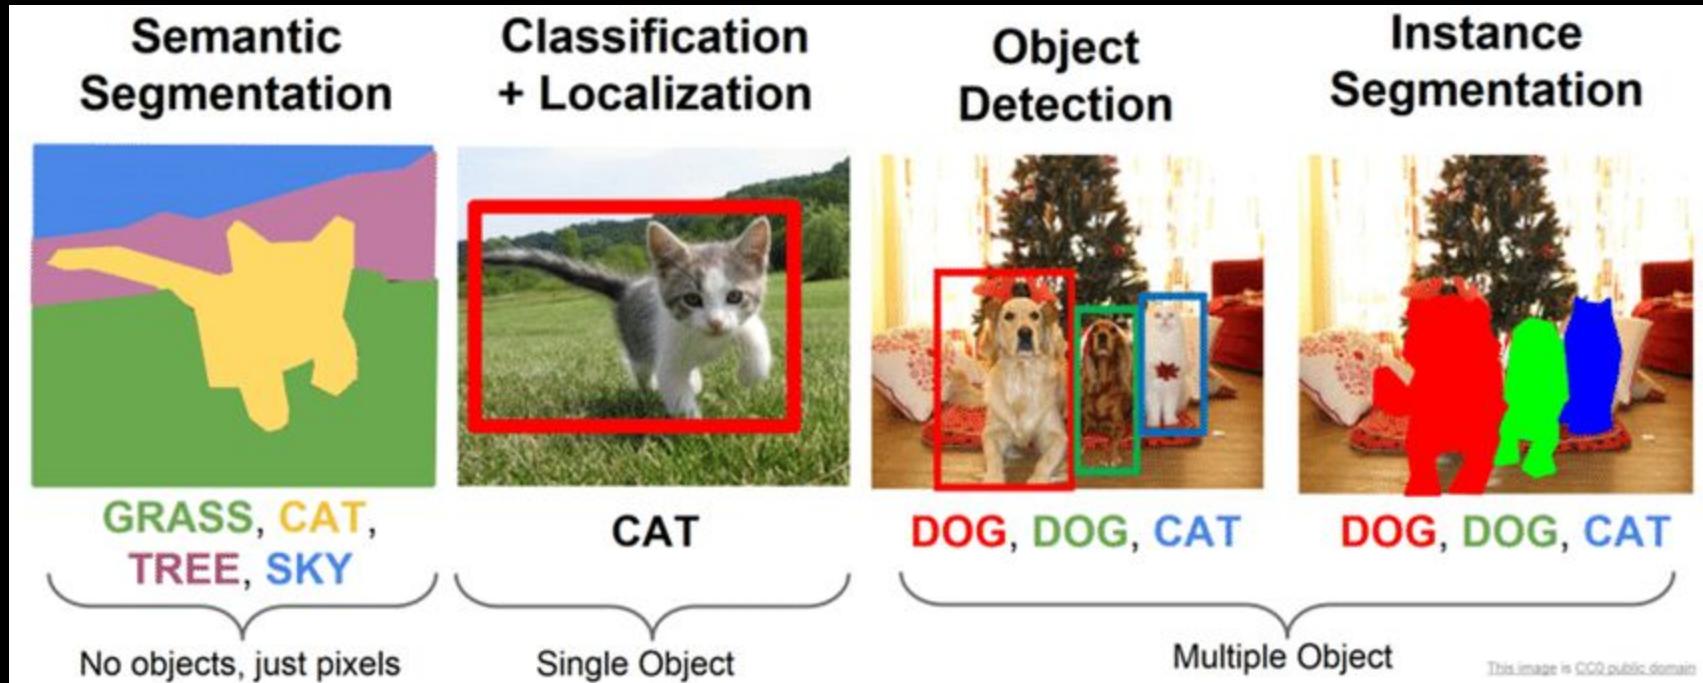
Image-wide pixel-to-pixel

Mostly used in biomedical imaging

Metrics: Jaccard index,  $F_1$ -score

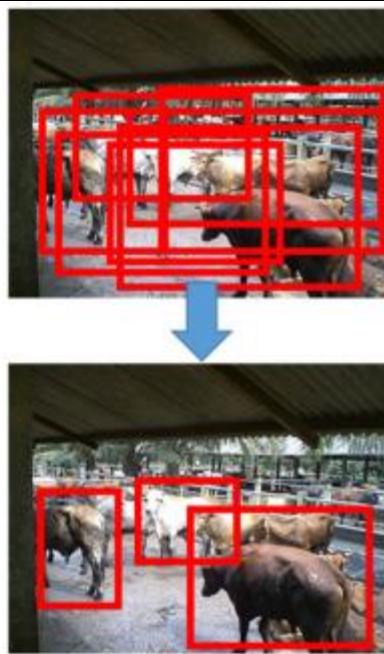
Usually pretrained on MS-COCO

*Deep Learning Computer Vision methods*  
— Segmentation types —



## *Deep Learning Computer Vision methods*

### — Segmentation Intermezzo: Non-Max Suppression —

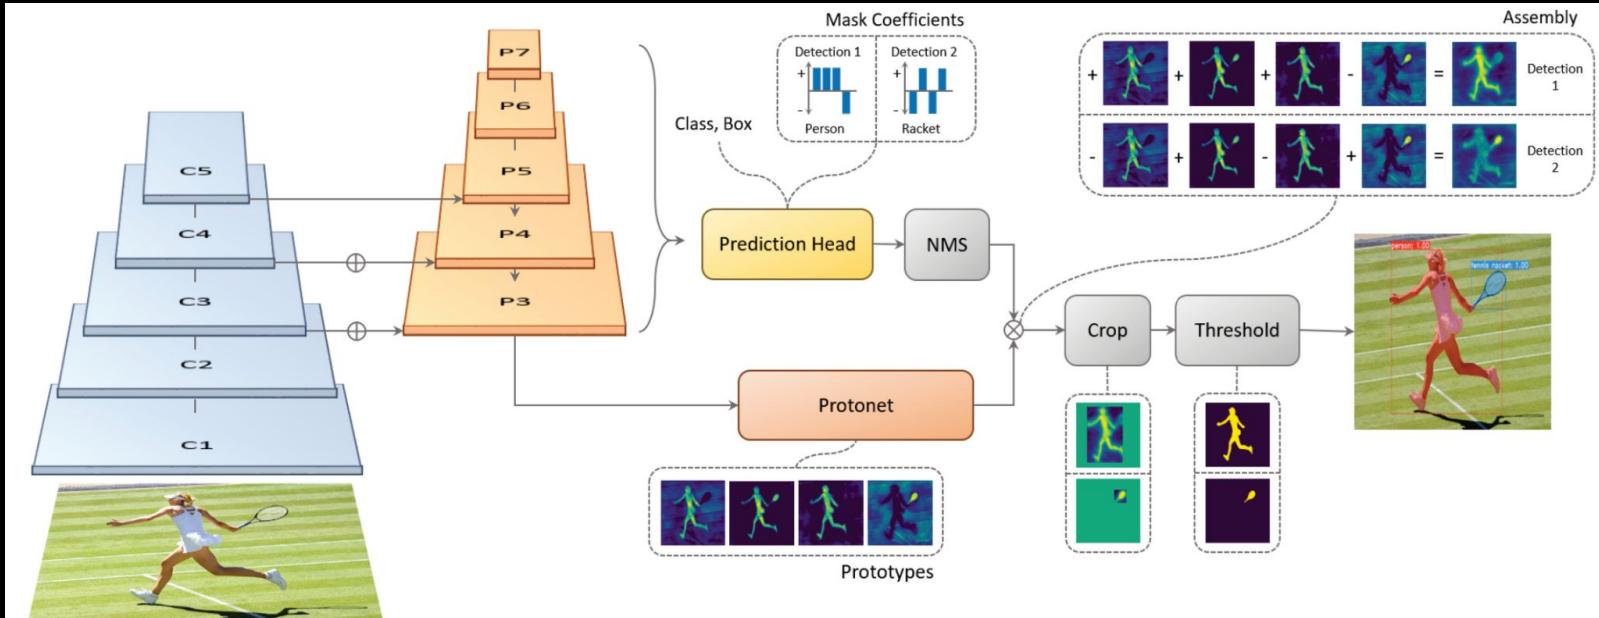


Algorithm:

1. Select the box  $s$  with highest confidence, remove it from  $P$ , and add it to  $K$ .
2. If  $\text{IoU}(s, p) > \theta$  for any  $p \in P$ , then remove  $p$  from  $P$ .
3. Repeat steps 1 & 2 until  $P$  is empty.
4. The list  $K$  contains the final boxes.

# Deep Learning Computer Vision methods

## — YOLACT —



**Instance segmentation:**

Image to (mask, class) pairs

Metrics: mAP,  $F_1$ -score

Insightful

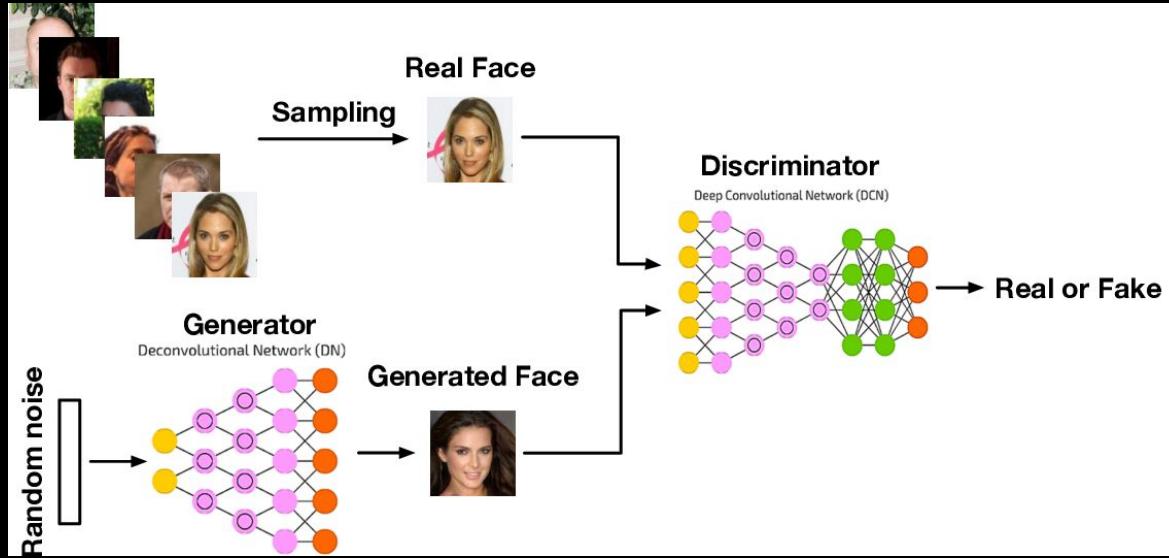
**Algorithm:**

1. Prototype generation branch (semantic segmentation, but supervised by heads)
2. Mask coefficient branch (object detection-like, then non-max suppression)
3. Mask assembly (Linear combination of prototypes with coefficients from masks)
4. Crop, sigmoid, threshold to produce result.



## Deep Learning Computer Vision methods

### — Generation: Adversarial networks —

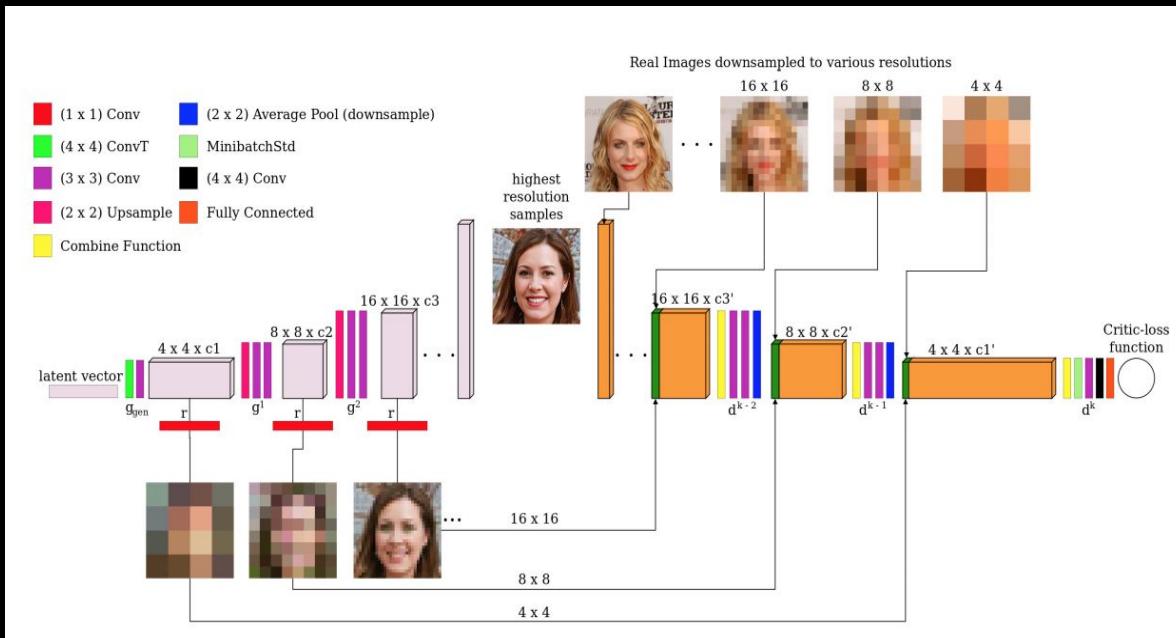


- Pitting two networks against each other
- Generator metric: *Fréchet inception distance*
  - Compares distributions of Inception V3 feature maps of real and generated images
- Common challenge: overpowered discriminators

# Deep Learning Computer Vision methods

## — Generation: StyleGAN2 —

- Smooth latent space between any two faces
  - Through complex normalization layers
- Skip connections to improve upsampling



# *Deep Learning Computer Vision methods*

## *— Generation: Neural style transfer —*

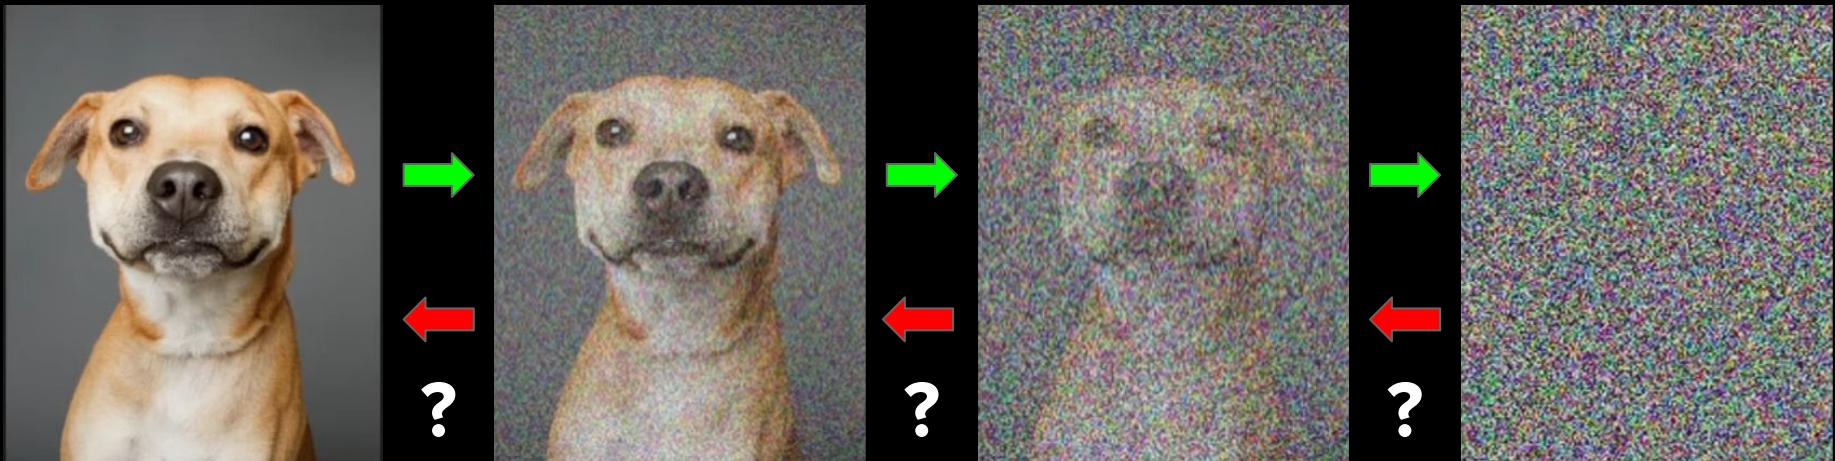
Use the smooth latent space to combine low-level feature maps (“style”) from one image, and the high-level feature maps (pose, facial structure) of another.



*Deep Learning Computer Vision methods*  
— Generation: Segmentation Masks (GauGAN) —



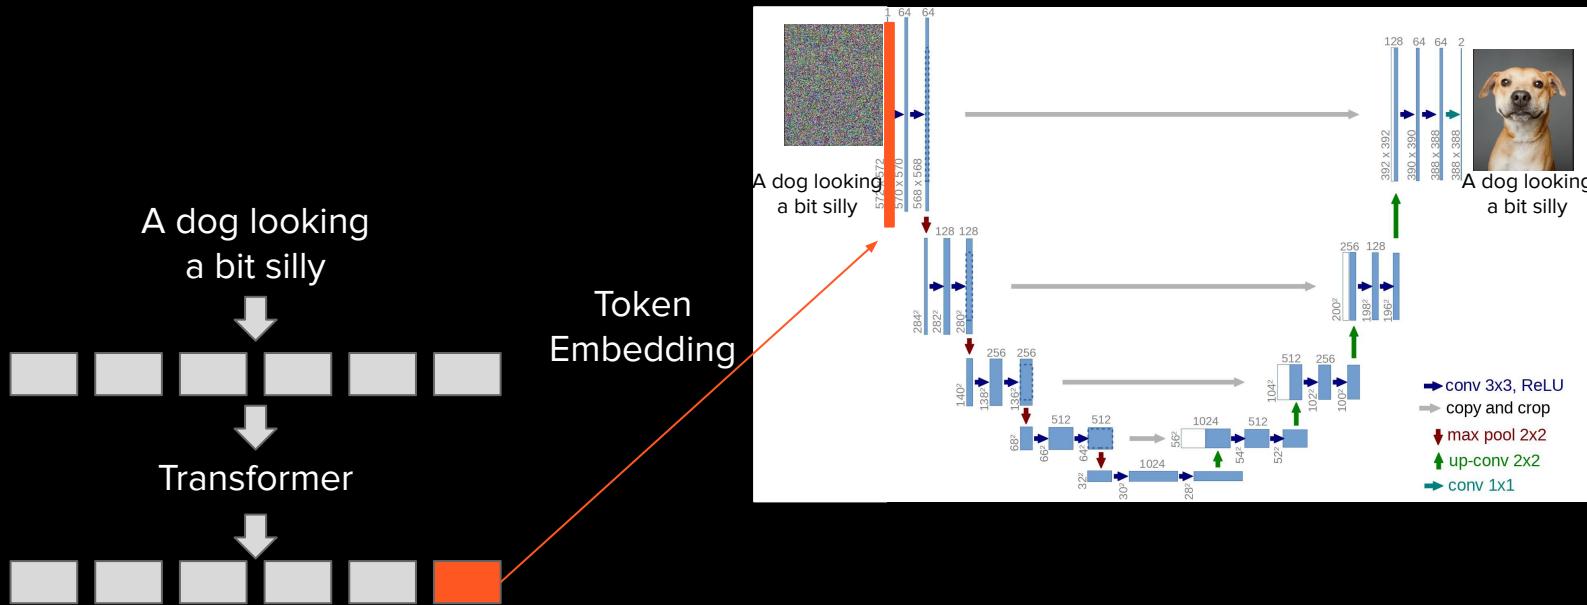
*Deep Learning Computer Vision methods*  
— Generation: Diffusion Models —



A dog looking  
a bit silly

# *Deep Learning Computer Vision methods*

## *— Generation: Diffusion Models —*



# *Deep Learning Computer Vision methods*

## — Generation: Diffusion Models —



panda mad scientist mixing sparkling chemicals, artstation



a corgi's head depicted as an explosion of a nebula



a propaganda poster depicting a cat dressed as french emperor napoleon holding a piece of cheese



a teddy bear on a skateboard in times square

DALL-E 2



Sprouts in the shape of text 'Imagen' coming out of a fairytale book.



A photo of a Shiba Inu dog with a backpack riding a bike. It is wearing sunglasses and a beach hat.



Teddy bears swimming at the Olympics 400m Butterfly event.



A cute corgi lives in a house made out of sushi.

IMAGEN



Diorama:2 wave carved from transparent glass,  
glowing inside, atmospheric, 3d render, hyperrealistic,  
redshift render, -q 2



The last supper, cyberpunk

MidJourney

## *Deep Learning Computer Vision methods — Generation: Diffusion Models —*



Cat Knight



Jaba The Trump

Stable Diffusion

“Artificial intelligence consultants getting ready for a game hackathon”

Go to: <https://github.com/Sam-Sweere-Vantage-AI/cv-hackathon>

